

Assignment 3 Solution

Table of Contents

[Problem 1. Steps](#)
[Problem 2. Steps](#)
[Problem 3. Steps](#)
[Problem 4. Steps](#)
[Problem 5. Steps](#)

[Problem 1. Results](#)
[Problem 2. Results](#)
[Problem 3 Results](#)
[Problem 4 Results](#)
[Problem 5 Results](#)

Problem 1) Modify attached class WordCount.java so that its result excludes the following stop words:

I		
a	in	
about	is	who
an	it	will
are	of	with
as	on	the
at	or	www
be	that	
by	the	
com	this	
for	to	
from	was	
how	what	
	when	
	where	

as well as special characters (dashes, parentases, etc). Stop word lists could be much longer than this. You do not have to be extremely thoroughful. You would like to get a more or less clean list of ordinary words with the numbers of their occurrences. Do not fret. Be reasonable. Perform analysis on the text of James Joyce's Ulysis.

Problem 1. Steps

Program overview: the stop words, upper or lower case, and any
words made up of special characters without any alphanumeric characters

are removed.

#Note: I renamed my class to **WordCountTwo.java** to prevent running the wrong # program. The first 20 results are listed in the appendix of this document.

First I exported the Hadoop classpath to a variable HCP to make the # compiling commands more convenient

```
[joe@localhost examples]$ export
HCP=/etc/hadoop/conf:/usr/lib/hadoop/lib/*:/usr/lib/hadoop/./*/:/usr/lib/hadoop-hdfs/./:/u
sr/lib/hadoop-hdfs/lib/*:/usr/lib/hadoop-hdfs/./*/:/usr/lib/hadoop-yarn/lib/*:/usr/lib/hadoo
p-yarn/./*/:/usr/lib/hadoop-mapreduce/lib/*:/usr/lib/hadoop-mapreduce/.//*
```

I compiled my program **WordCountTwo.java**

```
[joe@localhost examples]$ javac -classpath $HCP -d . WordCountTwo.java
```

I created a jar with my class files for WordCountTwo

```
[joe@localhost examples]$ jar -cvf wordcount.jar org/*
added manifest
adding: org/apache/(in = 0) (out= 0)(stored 0%)
adding: org/apache/hadoop/(in = 0) (out= 0)(stored 0%)
adding: org/apache/hadoop/examples/(in = 0) (out= 0)(stored 0%)
adding: org/apache/hadoop/examples/WordCountTwo$TokenizerMapper.class(in =
2636) (out= 1302)(deflated 50%)
adding: org/apache/hadoop/examples/WordCountTwo$IntSumReducer.class(in = 1802)
(out= 754)(deflated 58%)
adding: org/apache/hadoop/examples/WordCountTwo.class(in = 2005) (out=
1058)(deflated 47%)
```

Finally I successfully ran WordCountTwo on the Ulysses text

```
[joe@localhost examples]$ hadoop jar wordcount.jar
org.apache.hadoop.examples.WordCountTwo ulysses output
16/02/18 18:38:21 INFO client.RMProxy: Connecting to ResourceManager at
/0.0.0.0:8032
16/02/18 18:38:23 INFO input.FileInputFormat: Total input paths to process : 1
16/02/18 18:38:23 INFO mapreduce.JobSubmitter: number of splits:1
16/02/18 18:38:23 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_1455825540424_0009
16/02/18 18:38:24 INFO impl.YarnClientImpl: Submitted application
application_1455825540424_0009
16/02/18 18:38:24 INFO mapreduce.Job: The url to track the job:
http://localhost:8088/proxy/application_1455825540424_0009/
```

16/02/18 18:38:24 INFO mapreduce.Job: Running job: job_1455825540424_0009
16/02/18 18:38:37 INFO mapreduce.Job: Job job_1455825540424_0009 running in uber
mode : false
16/02/18 18:38:37 INFO mapreduce.Job: map 0% reduce 0%
16/02/18 18:38:51 INFO mapreduce.Job: map 100% reduce 0%
16/02/18 18:39:02 INFO mapreduce.Job: map 100% reduce 100%
16/02/18 18:39:03 INFO mapreduce.Job: Job job_1455825540424_0009 completed
successfully
16/02/18 18:39:03 INFO mapreduce.Job: Counters: 49

File System Counters

FILE: Number of bytes read=724071
FILE: Number of bytes written=1671633
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=1573191
HDFS: Number of bytes written=526820
HDFS: Number of read operations=6
HDFS: Number of large read operations=0
HDFS: Number of write operations=2

Job Counters

Launched map tasks=1
Launched reduce tasks=1
Data-local map tasks=1
Total time spent by all maps in occupied slots (ms)=12330
Total time spent by all reduces in occupied slots (ms)=7636
Total time spent by all map tasks (ms)=12330
Total time spent by all reduce tasks (ms)=7636
Total vcore-seconds taken by all map tasks=12330
Total vcore-seconds taken by all reduce tasks=7636
Total megabyte-seconds taken by all map tasks=12625920
Total megabyte-seconds taken by all reduce tasks=7819264

Map-Reduce Framework

Map input records=33056
Map output records=201025
Map output bytes=2103796
Map output materialized bytes=724071
Input split bytes=112
Combine input records=201025
Combine output records=49998
Reduce input groups=49998
Reduce shuffle bytes=724071
Reduce input records=49998
Reduce output records=49998
Spilled Records=99996

```
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=428
CPU time spent (ms)=7730
Physical memory (bytes) snapshot=322125824
Virtual memory (bytes) snapshot=5440196608
Total committed heap usage (bytes)=165810176

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=1573079
File Output Format Counters
  Bytes Written=526820
```

Problem 2) Write another MapReduce program which would read the “word count” output of the previous job and order the results by the declining number of occurrences.

Problem 2. Steps

Program overview: It sorts words by count into descending order. My program enters # the count in as the key in the map phase and uses a custom key sorter to get descending # order. It then writes out the count and the word in the reduce phase.

As in #1, I compiled my java program named **WordCountThree.java** for this # problem and created a jar with the class file named wordcount.jar

```
[joe@localhost examples]$ javac -classpath $HCP -d . WordCountThree.java
[joe@localhost examples]$ jar -cvf wordcount.jar org/*
added manifest
adding: org/apache/(in = 0) (out= 0)(stored 0%)
adding: org/apache/hadoop/(in = 0) (out= 0)(stored 0%)
adding: org/apache/hadoop/examples/(in = 0) (out= 0)(stored 0%)
adding: org/apache/hadoop/examples/WordCountThree$Reverse.class(in = 643) (out=
355)(deflated 44%)
adding: org/apache/hadoop/examples/WordCountThree.class(in = 2223) (out=
1138)(deflated 48%)
adding: org/apache/hadoop/examples/WordCountThree$TokenizerMapper.class(in =
1837) (out= 754)(deflated 58%)
```

adding: org/apache/hadoop/examples/WordCountThree\$IntSumReducer.class(in = 2040)
(out= 876)(deflated 57%)

I ran the program **WordCountThree** that reorders the words by number of
occurrences in descending order. The file with WordCountTwo's output is renamed
to 1044 and the output of this program is 1051. Additionally, the first 20 results
are listed in the appendix for references.

```
[joe@localhost examples]$ hadoop jar wordcount.jar  
org.apache.hadoop.examples.WordCountThree 1044 1051  
16/02/19 19:51:14 INFO client.RMProxy: Connecting to ResourceManager at  
/0.0.0.0:8032  
16/02/19 19:51:16 INFO input.FileInputFormat: Total input paths to process : 1  
16/02/19 19:51:16 INFO mapreduce.JobSubmitter: number of splits:1  
16/02/19 19:51:17 INFO mapreduce.JobSubmitter: Submitting tokens for job:  
job_1455938417648_0002  
16/02/19 19:51:17 INFO impl.YarnClientImpl: Submitted application  
application_1455938417648_0002  
16/02/19 19:51:17 INFO mapreduce.Job: The url to track the job:  
http://localhost:8088/proxy/application_1455938417648_0002/  
16/02/19 19:51:17 INFO mapreduce.Job: Running job: job_1455938417648_0002  
16/02/19 19:51:32 INFO mapreduce.Job: Job job_1455938417648_0002 running in uber  
mode : false  
16/02/19 19:51:32 INFO mapreduce.Job: map 0% reduce 0%  
16/02/19 19:51:44 INFO mapreduce.Job: map 100% reduce 0%  
16/02/19 19:51:57 INFO mapreduce.Job: map 100% reduce 100%  
16/02/19 19:51:58 INFO mapreduce.Job: Job job_1455938417648_0002 completed  
successfully  
16/02/19 19:51:58 INFO mapreduce.Job: Counters: 49  
    File System Counters  
        FILE: Number of bytes read=724071  
        FILE: Number of bytes written=1672701  
        FILE: Number of read operations=0  
        FILE: Number of large read operations=0  
        FILE: Number of write operations=0  
        HDFS: Number of bytes read=526933  
        HDFS: Number of bytes written=526820  
        HDFS: Number of read operations=6  
        HDFS: Number of large read operations=0  
        HDFS: Number of write operations=2  
    Job Counters  
        Launched map tasks=1  
        Launched reduce tasks=1  
        Data-local map tasks=1  
        Total time spent by all maps in occupied slots (ms)=10662
```

Total time spent by all reduces in occupied slots (ms)=9346
Total time spent by all map tasks (ms)=10662
Total time spent by all reduce tasks (ms)=9346
Total vcore-seconds taken by all map tasks=10662
Total vcore-seconds taken by all reduce tasks=9346
Total megabyte-seconds taken by all map tasks=10917888
Total megabyte-seconds taken by all reduce tasks=9570304

Map-Reduce Framework

Map input records=49998
Map output records=49998
Map output bytes=624069
Map output materialized bytes=724071
Input split bytes=113
Combine input records=49998
Combine output records=49998
Reduce input groups=236
Reduce shuffle bytes=724071
Reduce input records=49998
Reduce output records=49998
Spilled Records=99996
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=690
CPU time spent (ms)=5470
Physical memory (bytes) snapshot=321998848
Virtual memory (bytes) snapshot=5445128192
Total committed heap usage (bytes)=165810176

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=526820

File Output Format Counters

Bytes Written=526820

Problem 3) Create a program that will find out how many words appear only once, how many twice, three times, four times and so on in James Joyce's Ulysis

Problem 3. Steps

WordCountFour overview: It determines how many words occur once, twice, three
times, and so on. It enters a word's count as the key and an IntWritable value of one
during mapping. It then finds the sum of IntWritable's in the reducing phase for each
word count.

I compile my program named **WordCountFour**, put the class file into a jar.

```
[joe@localhost examples]$ javac -classpath $HCP -d . WordCountFour.java
[joe@localhost examples]$ jar -cvf wordcount.jar org/*added manifest
adding: org/apache/(in = 0) (out= 0)(stored 0%)
adding: org/apache/hadoop/(in = 0) (out= 0)(stored 0%)
adding: org/apache/hadoop/examples/(in = 0) (out= 0)(stored 0%)
adding: org/apache/hadoop/examples/WordCountFour.class(in = 2099) (out=
1088)(deflated 48%)
adding: org/apache/hadoop/examples/WordCountFour$TokenizerMapper.class(in =
1834) (out= 752)(deflated 58%)
adding: org/apache/hadoop/examples/WordCountFour$IntSumReducer.class(in = 2032)
(out= 846)(deflated 58%)
```

I run the mapreduce program with WordCountTwo's output as the input in the file
firstoutput and this program outputs to the folder named thirtyone

```
[joe@localhost examples]$ hadoop jar wordcount.jar
org.apache.hadoop.examples.WordCountFour firstoutput thirtyone
16/02/18 22:31:59 INFO client.RMProxy: Connecting to ResourceManager at
/0.0.0.0:8032
16/02/18 22:32:01 INFO input.FileInputFormat: Total input paths to process : 1
16/02/18 22:32:01 INFO mapreduce.JobSubmitter: number of splits:1
16/02/18 22:32:01 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_1455825540424_0031
16/02/18 22:32:02 INFO impl.YarnClientImpl: Submitted application
application_1455825540424_0031
16/02/18 22:32:02 INFO mapreduce.Job: The url to track the job:
http://localhost:8088/proxy/application_1455825540424_0031/
16/02/18 22:32:02 INFO mapreduce.Job: Running job: job_1455825540424_0031
16/02/18 22:32:15 INFO mapreduce.Job: Job job_1455825540424_0031 running in uber
mode : false
16/02/18 22:32:15 INFO mapreduce.Job: map 0% reduce 0%
16/02/18 22:32:25 INFO mapreduce.Job: map 100% reduce 0%
16/02/18 22:32:34 INFO mapreduce.Job: map 100% reduce 100%
16/02/18 22:32:35 INFO mapreduce.Job: Job job_1455825540424_0031 completed
successfully
16/02/18 22:32:36 INFO mapreduce.Job: Counters: 49
    File System Counters
        FILE: Number of bytes read=1966
```

FILE: Number of bytes written=228115
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=526937
HDFS: Number of bytes written=1316
HDFS: Number of read operations=6
HDFS: Number of large read operations=0
HDFS: Number of write operations=2

Job Counters

Launched map tasks=1
Launched reduce tasks=1
Data-local map tasks=1
Total time spent by all maps in occupied slots (ms)=7996
Total time spent by all reduces in occupied slots (ms)=6703
Total time spent by all map tasks (ms)=7996
Total time spent by all reduce tasks (ms)=6703
Total vcore-seconds taken by all map tasks=7996
Total vcore-seconds taken by all reduce tasks=6703
Total megabyte-seconds taken by all map tasks=8187904
Total megabyte-seconds taken by all reduce tasks=6863872

Map-Reduce Framework

Map input records=49998
Map output records=49998
Map output bytes=624069
Map output materialized bytes=1966
Input split bytes=117
Combine input records=49998
Combine output records=236
Reduce input groups=236
Reduce shuffle bytes=1966
Reduce input records=236
Reduce output records=236
Spilled Records=472
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=303
CPU time spent (ms)=3430
Physical memory (bytes) snapshot=314617856
Virtual memory (bytes) snapshot=5445595136
Total committed heap usage (bytes)=165810176

Shuffle Errors

BAD_ID=0
CONNECTION=0


```
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=526820
File Output Format Counters
  Bytes Written=1316
```

Problem 4) Combine operations of two MapReduce programs in Problems 1 and 3 above into a single program with chained MapReduce jobs.

Problem 4. Steps

WordCountChained overview: combine solutions from #1 and #3 with mapreduce chaining and a temporary storage file. Delete the temporary file after both jobs complete.

Note: First 20 results are listed in the appendix.

I compile and jar the class file for the class **WordCountChained**

```
[joe@localhost examples]$ javac -classpath $HCP -d . WordCountChained.java
[joe@localhost examples]$ jar -cvf wordcount.jar org/*
added manifest
adding: org/apache/(in = 0) (out= 0)(stored 0%)
adding: org/apache/hadoop/(in = 0) (out= 0)(stored 0%)
adding: org/apache/hadoop/examples/(in = 0) (out= 0)(stored 0%)
adding: org/apache/hadoop/examples/WordCountChained$TokenizerMapper.class(in =
1822) (out= 744)(deflated 59%)
adding: org/apache/hadoop/examples/WordCountChained$TokenizerMapperOne.class(in
= 2654) (out= 1318)(deflated 50%)
adding: org/apache/hadoop/examples/WordCountChained.class(in = 2895) (out=
1291)(deflated 55%)
adding: org/apache/hadoop/examples/WordCountChained$IntSumReducer.class(in =
1870) (out= 771)(deflated 58%)
adding: org/apache/hadoop/examples/WordCountChained$IntSumReducerOne.class(in =
1820) (out= 760)(deflated 58%)
```

I run the **WordCountChained** program on the ulysses file and the two mapreduce
processes output below. Also the output file is named 1109

```
[joe@localhost examples]$ hadoop jar wordcount.jar
org.apache.hadoop.examples.WordCountChained ulysses 1109
16/02/19 20:09:33 INFO client.RMProxy: Connecting to ResourceManager at
/0.0.0.0:8032
```

16/02/19 20:09:34 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.

16/02/19 20:09:35 INFO input.FileInputFormat: Total input paths to process : 1

16/02/19 20:09:35 INFO mapreduce.JobSubmitter: number of splits:1

16/02/19 20:09:35 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_1455938417648_0004

16/02/19 20:09:36 INFO impl.YarnClientImpl: Submitted application
application_1455938417648_0004

16/02/19 20:09:36 INFO mapreduce.Job: The url to track the job:

http://localhost:8088/proxy/application_1455938417648_0004/

16/02/19 20:09:36 INFO mapreduce.Job: Running job: job_1455938417648_0004

16/02/19 20:09:50 INFO mapreduce.Job: Job job_1455938417648_0004 running in uber
mode : false

16/02/19 20:09:50 INFO mapreduce.Job: map 0% reduce 0%

16/02/19 20:10:06 INFO mapreduce.Job: map 46% reduce 0%

16/02/19 20:10:09 INFO mapreduce.Job: map 67% reduce 0%

16/02/19 20:10:14 INFO mapreduce.Job: map 100% reduce 0%

16/02/19 20:10:26 INFO mapreduce.Job: map 100% reduce 100%

16/02/19 20:10:27 INFO mapreduce.Job: Job job_1455938417648_0004 completed
successfully

16/02/19 20:10:28 INFO mapreduce.Job: Counters: 49

File System Counters

FILE: Number of bytes read=724071

FILE: Number of bytes written=1671399

FILE: Number of read operations=0

FILE: Number of large read operations=0

FILE: Number of write operations=0

HDFS: Number of bytes read=1573191

HDFS: Number of bytes written=526820

HDFS: Number of read operations=6

HDFS: Number of large read operations=0

HDFS: Number of write operations=2

Job Counters

Launched map tasks=1

Launched reduce tasks=1

Data-local map tasks=1

Total time spent by all maps in occupied slots (ms)=22381

Total time spent by all reduces in occupied slots (ms)=9802

Total time spent by all map tasks (ms)=22381

Total time spent by all reduce tasks (ms)=9802

Total vcore-seconds taken by all map tasks=22381

Total vcore-seconds taken by all reduce tasks=9802

Total megabyte-seconds taken by all map tasks=22918144

Total megabyte-seconds taken by all reduce tasks=10037248

Map-Reduce Framework

Map input records=33056
Map output records=201025
Map output bytes=2103796
Map output materialized bytes=724071
Input split bytes=112
Combine input records=201025
Combine output records=49998
Reduce input groups=49998
Reduce shuffle bytes=724071
Reduce input records=49998
Reduce output records=49998
Spilled Records=99996
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=521
CPU time spent (ms)=12040
Physical memory (bytes) snapshot=321773568
Virtual memory (bytes) snapshot=5440188416
Total committed heap usage (bytes)=165810176

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=1573079

File Output Format Counters

Bytes Written=526820

16/02/19 20:10:28 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032

16/02/19 20:10:28 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.

16/02/19 20:10:28 INFO input.FileInputFormat: Total input paths to process : 1

16/02/19 20:10:28 INFO mapreduce.JobSubmitter: number of splits:1

16/02/19 20:10:29 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1455938417648_0005

16/02/19 20:10:29 INFO impl.YarnClientImpl: Submitted application application_1455938417648_0005

16/02/19 20:10:29 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1455938417648_0005/

16/02/19 20:10:29 INFO mapreduce.Job: Running job: job_1455938417648_0005
16/02/19 20:10:44 INFO mapreduce.Job: Job job_1455938417648_0005 running in uber
mode : false
16/02/19 20:10:44 INFO mapreduce.Job: map 0% reduce 0%
16/02/19 20:11:03 INFO mapreduce.Job: map 100% reduce 0%
16/02/19 20:11:15 INFO mapreduce.Job: map 100% reduce 100%
16/02/19 20:11:16 INFO mapreduce.Job: Job job_1455938417648_0005 completed
successfully
16/02/19 20:11:16 INFO mapreduce.Job: Counters: 49

File System Counters

FILE: Number of bytes read=2366
FILE: Number of bytes written=227967
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=526939
HDFS: Number of bytes written=1388
HDFS: Number of read operations=6
HDFS: Number of large read operations=0
HDFS: Number of write operations=2

Job Counters

Launched map tasks=1
Launched reduce tasks=1
Data-local map tasks=1
Total time spent by all maps in occupied slots (ms)=17373
Total time spent by all reduces in occupied slots (ms)=9578
Total time spent by all map tasks (ms)=17373
Total time spent by all reduce tasks (ms)=9578
Total vcore-seconds taken by all map tasks=17373
Total vcore-seconds taken by all reduce tasks=9578
Total megabyte-seconds taken by all map tasks=17789952
Total megabyte-seconds taken by all reduce tasks=9807872

Map-Reduce Framework

Map input records=49998
Map output records=49998
Map output bytes=399984
Map output materialized bytes=2366
Input split bytes=119
Combine input records=49998
Combine output records=236
Reduce input groups=236
Reduce shuffle bytes=2366
Reduce input records=236
Reduce output records=236
Spilled Records=472

```
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=303
CPU time spent (ms)=8920
Physical memory (bytes) snapshot=316198912
Virtual memory (bytes) snapshot=5440196608
Total committed heap usage (bytes)=165810176
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=526820
File Output Format Counters
  Bytes Written=1388
```

Problem 5) Move attached `Inverter.java` class from old MapReduce API to the new API. Demonstrate that new and old class produce the same result. Use patent data set to demonstrate your work.

Problem 5. Steps

```
# I compile the InverterNewAPI.java into a class and create a jar
# for the class file. The jar is named wordcount.jar which is un descriptive
# but will function to run the program.
```

```
[joe@localhost inversion]$ javac -classpath $HCP -d . InverterNewAPI.java
[joe@localhost inversion]$ jar -cvf wordcount.jar edu/*
added manifest
adding: edu/hu/(in = 0) (out= 0)(stored 0%)
adding: edu/hu/bigdata/(in = 0) (out= 0)(stored 0%)
adding: edu/hu/bigdata/InverterNewAPI.class(in = 2352) (out= 1106)(deflated 52%)
adding: edu/hu/bigdata/InverterNewAPI$Reduce.class(in = 1868) (out= 805)(deflated 56%)
adding: edu/hu/bigdata/InverterNewAPI$MapClass.class(in = 1274) (out= 486)(deflated 61%)
```

```
# I run the InverterNewAPI program in hadoop. The input is the patents data and
# the output is a folder named 152.
```

```

[joe@localhost inversion]$ hadoop jar wordcount.jar edu.hu.bigdata.InverterNewAPI
patents 152
16/02/20 10:52:53 INFO client.RMProxy: Connecting to ResourceManager at
/0.0.0.0:8032
16/02/20 10:52:57 INFO input.FileInputFormat: Total input paths to process : 1
16/02/20 10:52:58 INFO mapreduce.JobSubmitter: number of splits:2
16/02/20 10:52:58 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_1455992537384_0003
16/02/20 10:52:59 INFO impl.YarnClientImpl: Submitted application
application_1455992537384_0003
16/02/20 10:52:59 INFO mapreduce.Job: The url to track the job:
http://localhost:8088/proxy/application_1455992537384_0003/
16/02/20 10:52:59 INFO mapreduce.Job: Running job: job_1455992537384_0003
16/02/20 10:53:12 INFO mapreduce.Job: Job job_1455992537384_0003 running in uber
mode : false
16/02/20 10:53:12 INFO mapreduce.Job: map 0% reduce 0%
16/02/20 10:53:33 INFO mapreduce.Job: map 8% reduce 0%
16/02/20 10:53:36 INFO mapreduce.Job: map 22% reduce 0%
16/02/20 10:53:40 INFO mapreduce.Job: map 26% reduce 0%
16/02/20 10:54:05 INFO mapreduce.Job: map 40% reduce 0%
16/02/20 10:54:08 INFO mapreduce.Job: map 47% reduce 0%
16/02/20 10:54:11 INFO mapreduce.Job: map 48% reduce 0%
16/02/20 10:54:29 INFO mapreduce.Job: map 56% reduce 0%
16/02/20 10:54:32 INFO mapreduce.Job: map 67% reduce 0%
16/02/20 10:54:54 INFO mapreduce.Job: map 68% reduce 0%
16/02/20 10:54:57 INFO mapreduce.Job: map 71% reduce 0%
16/02/20 10:55:00 INFO mapreduce.Job: map 75% reduce 0%
16/02/20 10:55:03 INFO mapreduce.Job: map 78% reduce 0%
16/02/20 10:55:06 INFO mapreduce.Job: map 81% reduce 0%
16/02/20 10:55:09 INFO mapreduce.Job: map 84% reduce 0%
16/02/20 10:55:12 INFO mapreduce.Job: map 88% reduce 0%
16/02/20 10:55:15 INFO mapreduce.Job: map 91% reduce 0%
16/02/20 10:55:18 INFO mapreduce.Job: map 94% reduce 0%
16/02/20 10:55:28 INFO mapreduce.Job: map 97% reduce 0%
16/02/20 10:55:31 INFO mapreduce.Job: map 100% reduce 0%
16/02/20 10:56:15 INFO mapreduce.Job: map 100% reduce 17%
16/02/20 10:56:21 INFO mapreduce.Job: map 100% reduce 67%
16/02/20 10:56:24 INFO mapreduce.Job: map 100% reduce 70%
16/02/20 10:56:27 INFO mapreduce.Job: map 100% reduce 73%
16/02/20 10:56:30 INFO mapreduce.Job: map 100% reduce 76%
16/02/20 10:56:33 INFO mapreduce.Job: map 100% reduce 79%
16/02/20 10:56:36 INFO mapreduce.Job: map 100% reduce 82%
16/02/20 10:56:39 INFO mapreduce.Job: map 100% reduce 84%
16/02/20 10:56:42 INFO mapreduce.Job: map 100% reduce 87%
16/02/20 10:56:45 INFO mapreduce.Job: map 100% reduce 90%

```

16/02/20 10:56:49 INFO mapreduce.Job: map 100% reduce 92%
16/02/20 10:56:52 INFO mapreduce.Job: map 100% reduce 94%
16/02/20 10:56:55 INFO mapreduce.Job: map 100% reduce 97%
16/02/20 10:56:58 INFO mapreduce.Job: map 100% reduce 100%
16/02/20 10:57:00 INFO mapreduce.Job: Job job_1455992537384_0003 completed successfully
16/02/20 10:57:00 INFO mapreduce.Job: Counters: 49

File System Counters

- FILE: Number of bytes read=594240678
- FILE: Number of bytes written=894018537
- FILE: Number of read operations=0
- FILE: Number of large read operations=0
- FILE: Number of write operations=0
- HDFS: Number of bytes read=264079761
- HDFS: Number of bytes written=158078539
- HDFS: Number of read operations=9
- HDFS: Number of large read operations=0
- HDFS: Number of write operations=2

Job Counters

- Launched map tasks=2
- Launched reduce tasks=1
- Data-local map tasks=2
- Total time spent by all maps in occupied slots (ms)=270943
- Total time spent by all reduces in occupied slots (ms)=86521
- Total time spent by all map tasks (ms)=270943
- Total time spent by all reduce tasks (ms)=86521
- Total vcore-seconds taken by all map tasks=270943
- Total vcore-seconds taken by all reduce tasks=86521
- Total megabyte-seconds taken by all map tasks=277445632
- Total megabyte-seconds taken by all reduce tasks=88597504

Map-Reduce Framework

- Map input records=16522439
- Map output records=16522439
- Map output bytes=264075431
- Map output materialized bytes=297120321
- Input split bytes=234
- Combine input records=0
- Combine output records=0
- Reduce input groups=3258984
- Reduce shuffle bytes=297120321
- Reduce input records=16522439
- Reduce output records=3258984
- Spilled Records=49567317
- Shuffled Maps =2
- Failed Shuffles=0

Merged Map outputs=2
GC time elapsed (ms)=5097
CPU time spent (ms)=140720
Physical memory (bytes) snapshot=500285440
Virtual memory (bytes) snapshot=8162627584
Total committed heap usage (bytes)=301146112
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
Bytes Read=264079527
File Output Format Counters
Bytes Written=158078539

I show the **InverterNewAPI** and **Inverter** programs both produce the same result in
the section with [problem #5 results](#). This is located at the end of this document,
[linked here](#).

If you have your own working VM with installed CDH5.5.1, do this assignment on that VM. If, for what ever reason, you do not posses a working VM with CDH5.5.1, please be free to download Clouder's Getting started VM and do you assignment on that VM.

Capture all steps of your implementation with comments indicating what is it you are accomplishing with every step in an MS Word document. Upload to the class site. Please upload your working Java files as well. Please post comments and questions to the class Discussion Board on the Canvas site.

APPENDIX

Problem 1. Results

Problem 1, First 20 Results

"Come 1
"Defects," 1
"I 1
"Information 1
"J" 1
"Plain 2
"Project 5
"Right 1
"Viator" 1
#4300] 1
\$5,000) 1
&c, 2
&c. 1
'46. 1
'92 1
'AS-IS' 1
'Slife, 1
'TWAS 1
'Tis 8
'Tis, 1

Problem 2. Results

Problem 2, First 20 Results:

6542 and
3035 his
2712 he
1505 her
1362 you
1113 him
1042 all
908 He
769 she
768 they
766 had
760 out
737 not
708 my
699 Mr
677 their
661 up
649 like
640 me
617 have

Problem 3 Results

Problem 3, First 20 Results

Note: Format is Occurences, Number of words that occur that many times

1	32735
2	7065
3	2932
4	1647
5	1102
6	727
7	509
8	409
9	337
10	261
11	190
12	165
13	130
14	124
15	112
16	106
17	83
18	85
19	64
20	65

Problem 4 Results

Problem 4, First 20 Results

1	32735
2	7065
3	2932
4	1647
5	1102
6	727
7	509
8	409
9	337
10	261
11	190
12	165
13	130
14	124
15	112
16	106

17	83
18	85
19	64
20	65

Problem 5 Results

First 20 Results from InverterNewAPI.java

"CITED"	"CITING"
1	3964859,4647229
10000	4539112
100000	5031388
1000006	4714284
1000007	4766693
1000011	5033339
1000017	3908629
1000026	4043055
1000033	4190903,4975983
1000043	4091523
1000044	4082383,4055371
1000045	4290571
1000046	5525001,5918892
1000049	5996916
1000051	4541310
1000054	4946631
1000065	4748968
1000067	5071294,4944640,5312208
1000070	4928425,5009029
1000073	4107819,5474494

20 Results from Inverter.java on the patents data

"CITED"	"CITING"
1	3964859,4647229
10000	4539112
100000	5031388
1000006	4714284
1000007	4766693
1000011	5033339
1000017	3908629
1000026	4043055
1000033	4190903,4975983
1000043	4091523
1000044	4082383,4055371
1000045	4290571

1000046	5525001,5918892
1000049	5996916
1000051	4541310
1000054	4946631
1000065	4748968
1000067	5071294,4944640,5312208
1000070	4928425,5009029
1000073	4107819,5474494

The previous two results, and the entire list of results, match proving that InverterNewAPI is updated to the new API correctly.

Additionally, below are the last 10 results of each to further show the results for the new and old API match.

InverterNewAPI, last 10 results

999961	5782495,5738381,5878901,4171117,4262874,5048788,4871140,4832301,4437639
999965	5052613
999968	3916735
999971	3965843
999972	4038129
999973	4900344,5427610
999974	5464105,4560073,4728158
999977	4092587
999978	3915443
999983	5143114,5394715,5806555

Inverter, last 10 results, which match the above results

999961	5782495,5738381,5878901,4171117,4262874,5048788,4871140,4832301,4437639
999965	5052613
999968	3916735
999971	3965843
999972	4038129
999973	4900344,5427610
999974	5464105,4560073,4728158
999977	4092587
999978	3915443
999983	5143114,5394715,5806555